

Achieving Near-Zero Hallucination In Large Language Models

Zoltán Blahovics Bence Nagy Krisztián Nemes-Kovács Balázs Fekete



Table of Contents

Introduction

Methodology

Results

Limitations and Further Research

Introduction

Methodology

- ▶ Experimental investigation in two phases:
 1. Baseline Establishment and Data Quality Analysis
 2. Architectural Intervention and Evaluation

Methodology - Phase 1

- ▶ Establishing model architecture and configuration for comparative experiments
 - ▶ Canonical transformer model
 - ▶ 6 encoder layers, 6 decoder layers, 8 attention heads
 - ▶ Standard parameter initialization
 - ▶ Adam optimizer with inverse square root scheduler

Methodology - Phase 1

- ▶ Creating two large training sets from publicly available datasets in the following steps:
 1. Classifying the data based on empirical reliability into categories:
 - ▶ High Reliability - exhibiting high factual consistency and validation
 - ▶ Low Reliability - exhibiting high factual heterogeneity and weak source attribution
 2. Based on the classification, formalizing the two distinct corpora:
 - ▶ Low-Reliability Corpus (\mathcal{D}_{LR})
 - ▶ High-Reliability Corpus (\mathcal{D}_{HR})
 3. Strict deduplication and factuality validation for \mathcal{D}_{HR}

Methodology - Phase 1

- ▶ Creating and training two baseline models
- ▶ Goal: isolate the effect of training data reliability on factual consistency
- ▶ Models: Baseline B_1 and Baseline B_2 were setup following to the previously demonstrated configuration
- ▶ Training B_1 and B_2 followed the identical training regimen, but used different data:
 - ▶ B_1 was trained exclusively on the \mathcal{D}_{LR}
 - ▶ B_2 was trained exclusively on the \mathcal{D}_{HR}

Methodology - Phase 1

- ▶ For measuring the effect of training data on factual adherence, a robust testing regimen was established.
- ▶ Metric: Hallucination Rate defined as the proportion of responses with ≥ 1 factually incorrect claim.
- ▶ $\mathcal{T}_{\text{Fact}}$ was created from evaluation prompts sampled equally from the *FEVER* dataset and *TruthfulQA* benchmark.
- ▶ $\mathcal{T}_{\text{Fact}}$ is used exclusively for internal benchmarking to measure improvement across development stages.
- ▶ Will be reused for measuring the improvement gained by data quality improvement (and for testing the novel architecture in a later step)

Methodology - Phase 2

- ▶ Novel architecture proposal
- ▶ The Layer-Specific Factual Gate (LSFG)
- ▶ Replaced the Standard Feed-Forward Networks (FFN) in the final decoder layers with a novel Gated Factual Network (GFN)
- ▶ The gating mechanism enforces selective suppression of activations contributing to factual inconsistency in the terminal layers.
- ▶ Formalization:
$$\text{Output} = g \odot \text{FFN}(z), \quad \text{where } g = \sigma(W_g z + b_g)$$
- ▶ The experimental model M_{LSFG} trained exclusively on \mathcal{D}_{HR}
- ▶ The internal benchmarking on $\mathcal{T}_{\text{Fact}}$ showed that both the data quality and the model architectural alterations contributed greatly to the reduction of hallucination rate.

Results

Limitations and Further Research

Bibliography

-  **Kyungjin Yoo and Rajeev Barua.** "Recovery of Object Oriented Features from C++ Binaries". *21st Asia-Pacific Software Engineering Conference*. Vol. 1. 231–238, (2014)
-  **Andre Pawlowski et al.** "MARX: Uncovering Class Hierarchies in C++ Programs"., *Network and Distributed System Security Symposium* (2017)
-  **Edward J. Schwartz et al.** "Using Logic Programming to Recover C++ Classes and Methods from Compiled Executables"., *ACM SIGSAC Conference on Computer and Communications Security*, 426–441, (2018)
-  **Rukayat Ayomide Erinfolami and Aravind Prakash..** "DeClassifier: Class-Inheritance Inference Engine for Optimized C++ Binaries", *arXiv eprints arXiv:1901.10073* (2019)